Introduction to data-parallelism
and OpenCL

Sylvain Lefebvre - INRIA

Quick reminders

Coalesced accesses

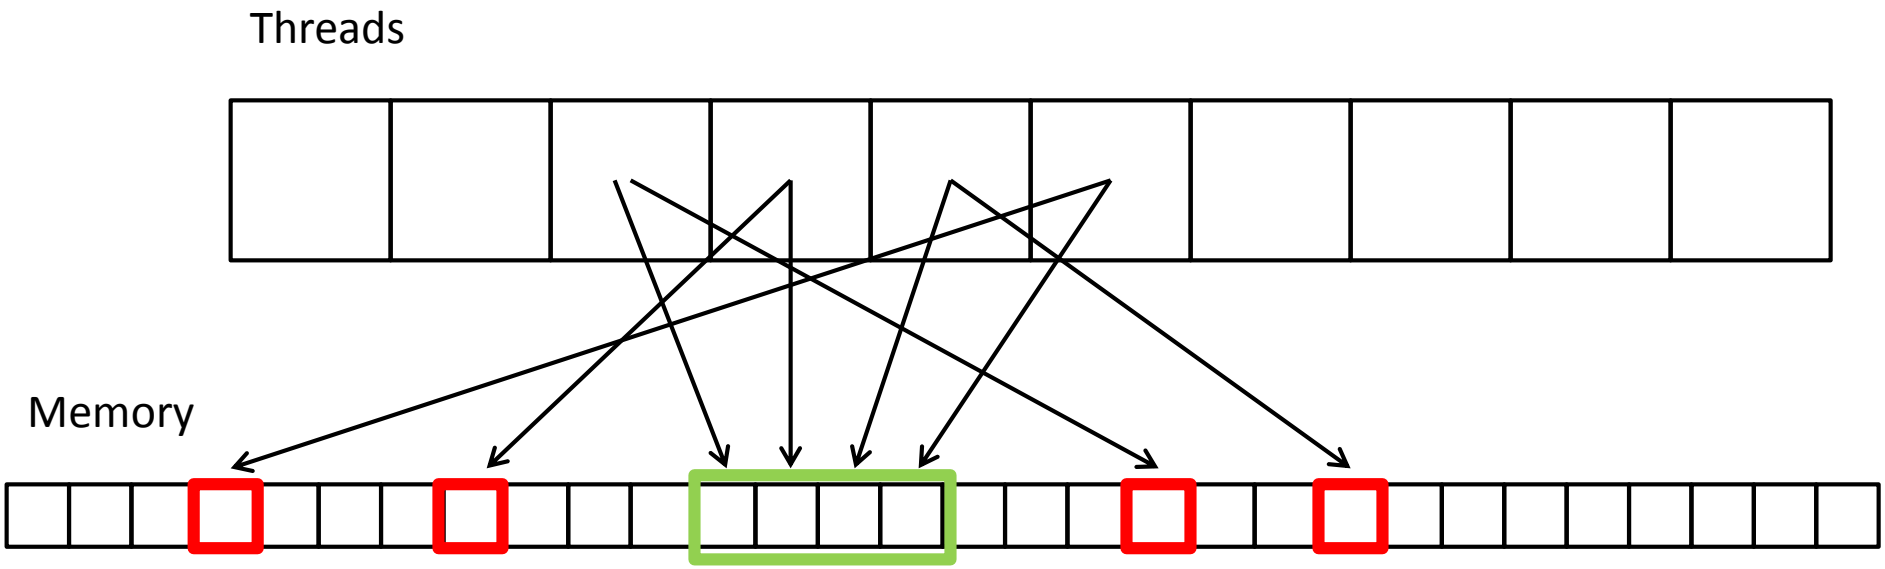

Pre-fetch

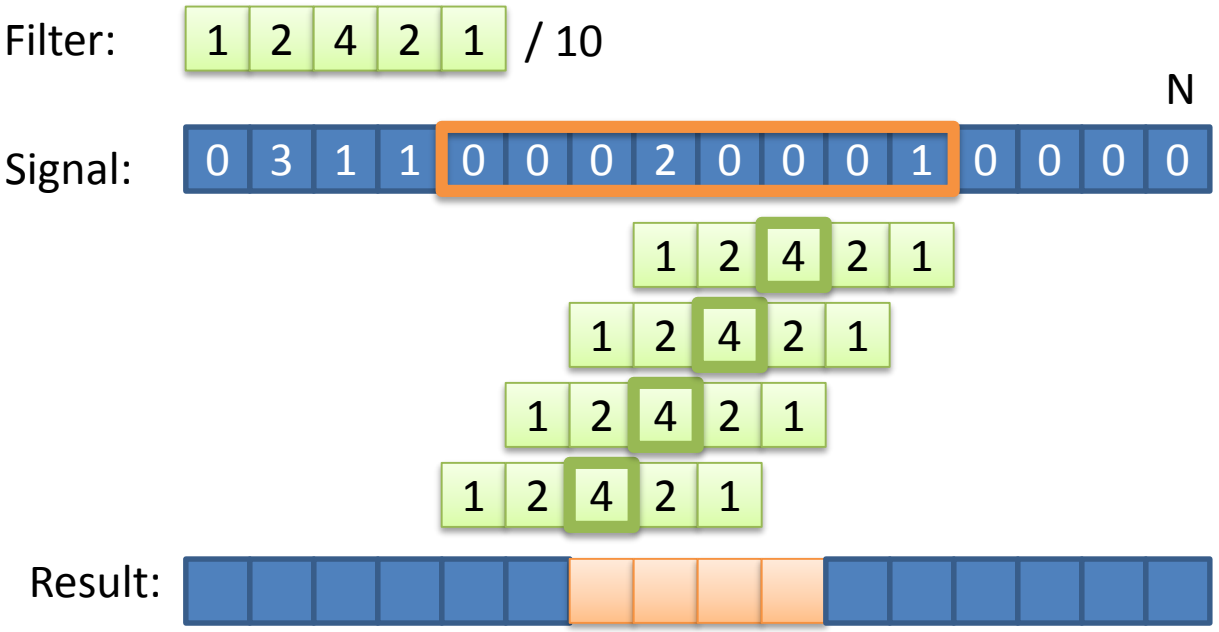
Filter:
1 2 4 2 1
/ 10
N
Signal:
0 3 1 1 0 0 0 2 0 0 0 1 0 0 0 0
1 2 4 2 1
1 2 4 2 1
1 2 4 2 1
1 2 4 2 1
Result:

Synchronization

- barrier
  - CLK_**LOCAL**_MEM_FANCE
  - CLK_**GLOBAL**_MEM_FANCE

- Atomics
  - inc/dec, add/sub, min/max, xchg/cmpxchg

Today

- Parallel Reduction
- Parallel Scan
- Applications

Parallel reduction

- Ex: Sum all entries of an array
- Sum-of-all is a special case of a *reduction*
- ‘Add’ could be replaced by
  - min / max / and / or / etc.
- Very important compute primitive
  - People have been working hard to optimize it

Reduction

- Assume ‘sum’ as operator
- First approach:
  - Atomic add
- Second approach:
  - Hierarchical sum

atomicAdd

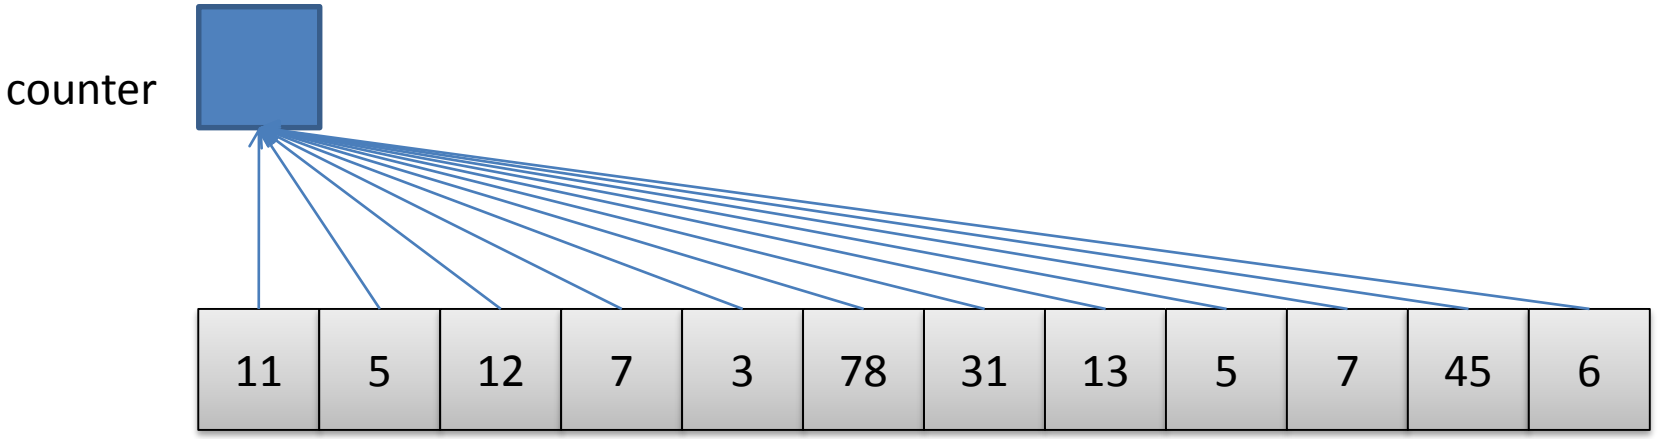

atomicAdd (2)

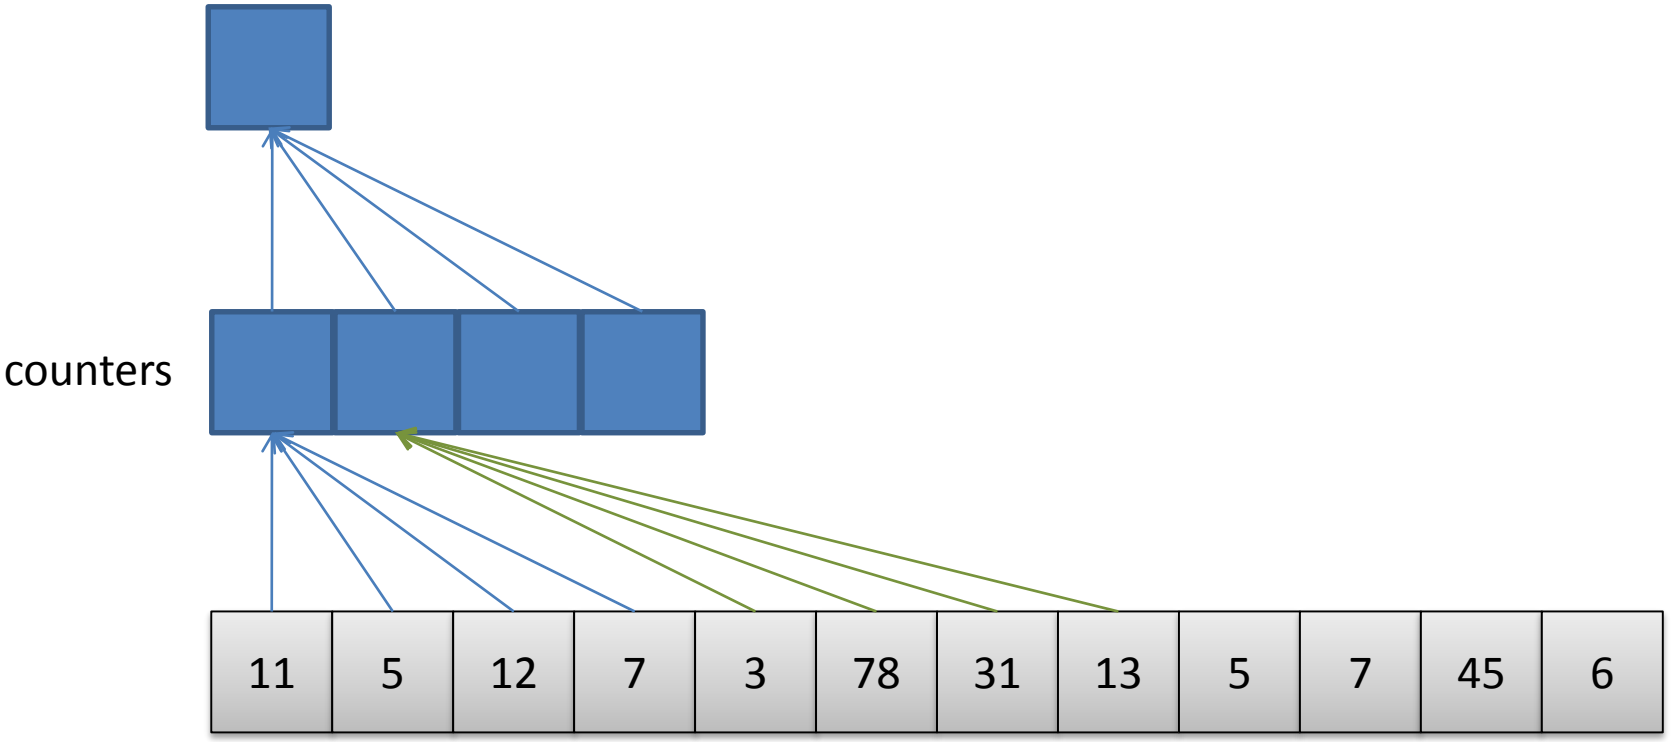

Hierarchy

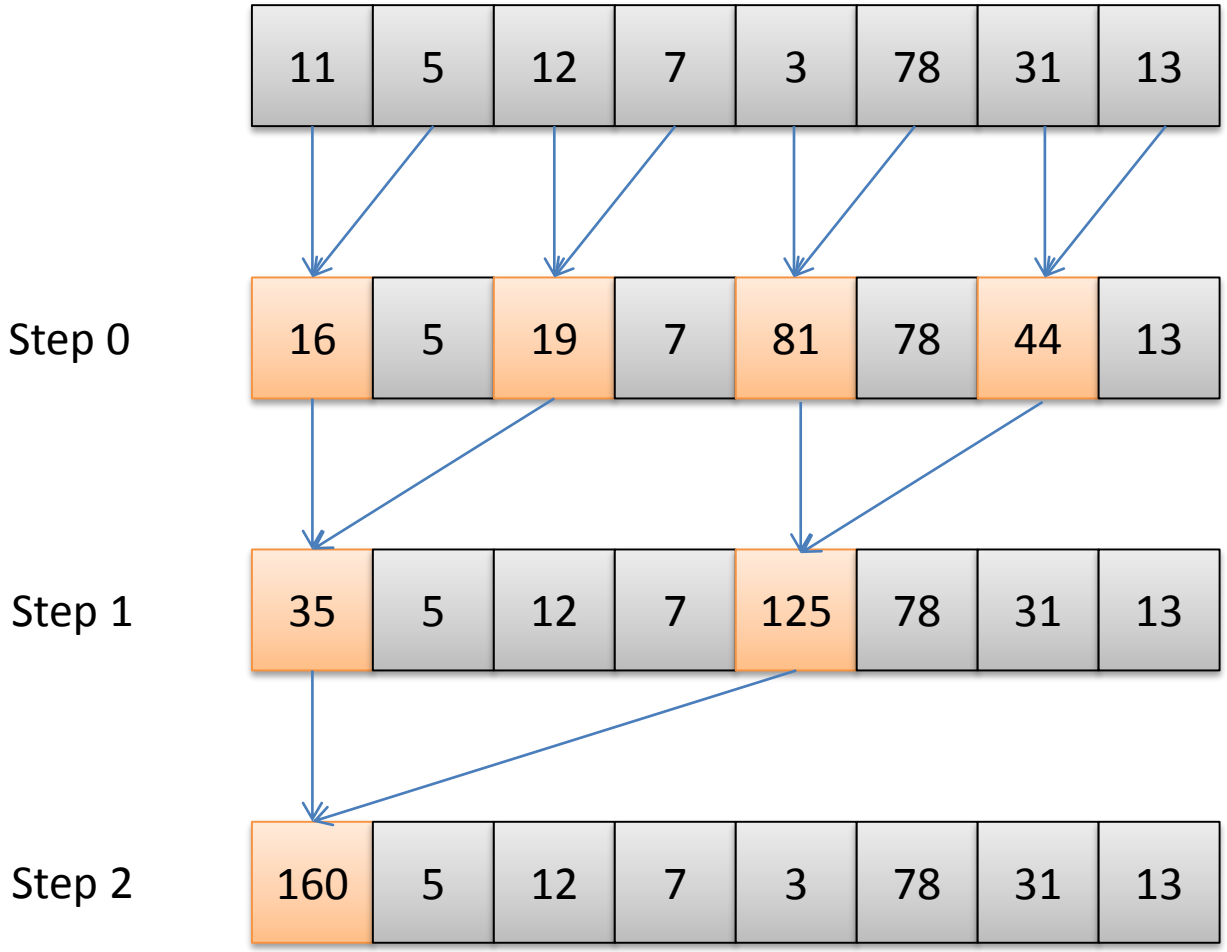
11 5 12 7 3 78 31 13
Step 0
16 5 19 7 81 78 44 13
Step 1
35 5 12 7 125 78 31 13
Step 2
160 5 12 7 3 78 31 13

Better

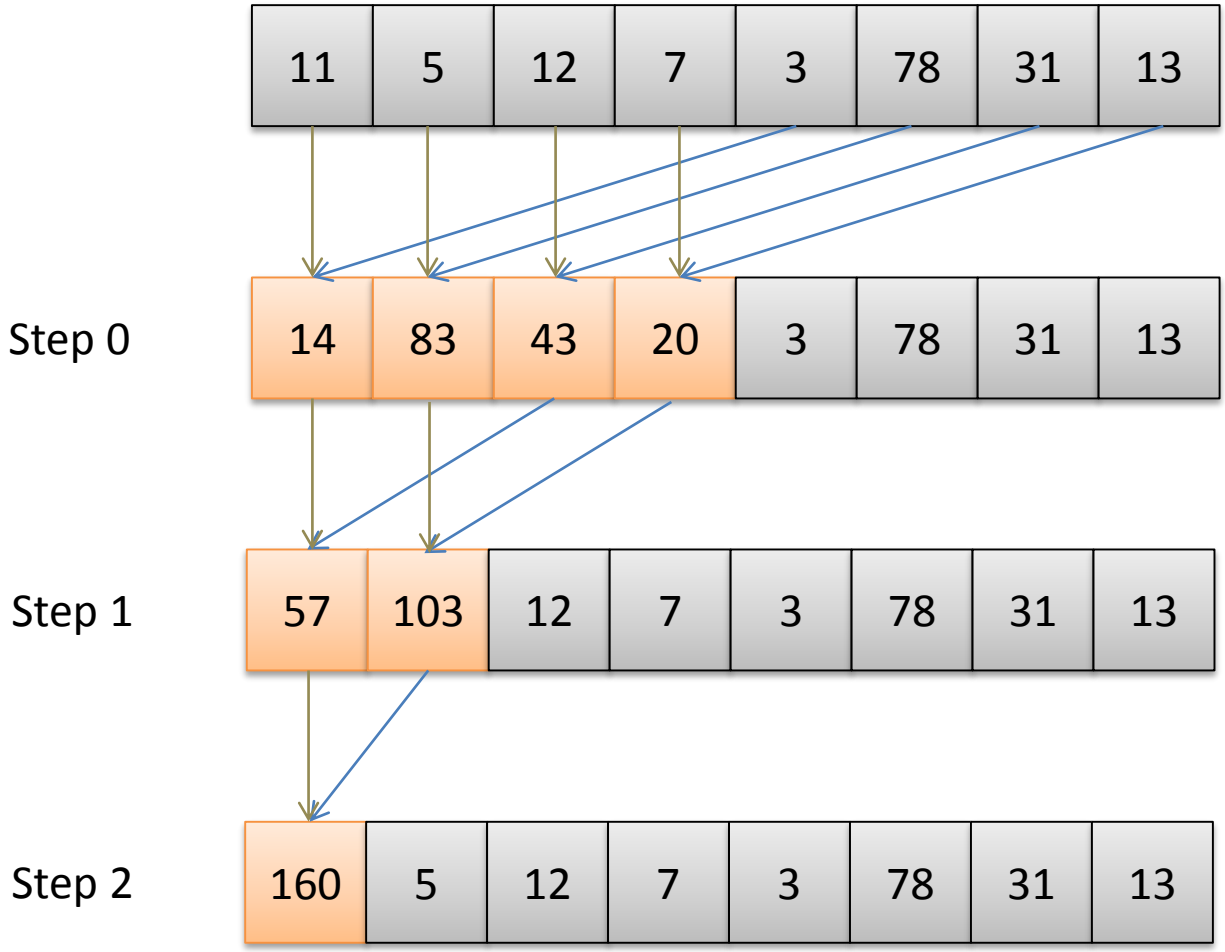
11 5 12 7 3 78 31 13
Step 0
14 83 43 20 3 78 31 13
Step 1
57 103 12 7 3 78 31 13
Step 2
160 5 12 7 3 78 31 13

Local vs. global

- Local is faster
- Limited due to synchronization

➔ Do local reduction on sub-arrays

Use local to minimize global

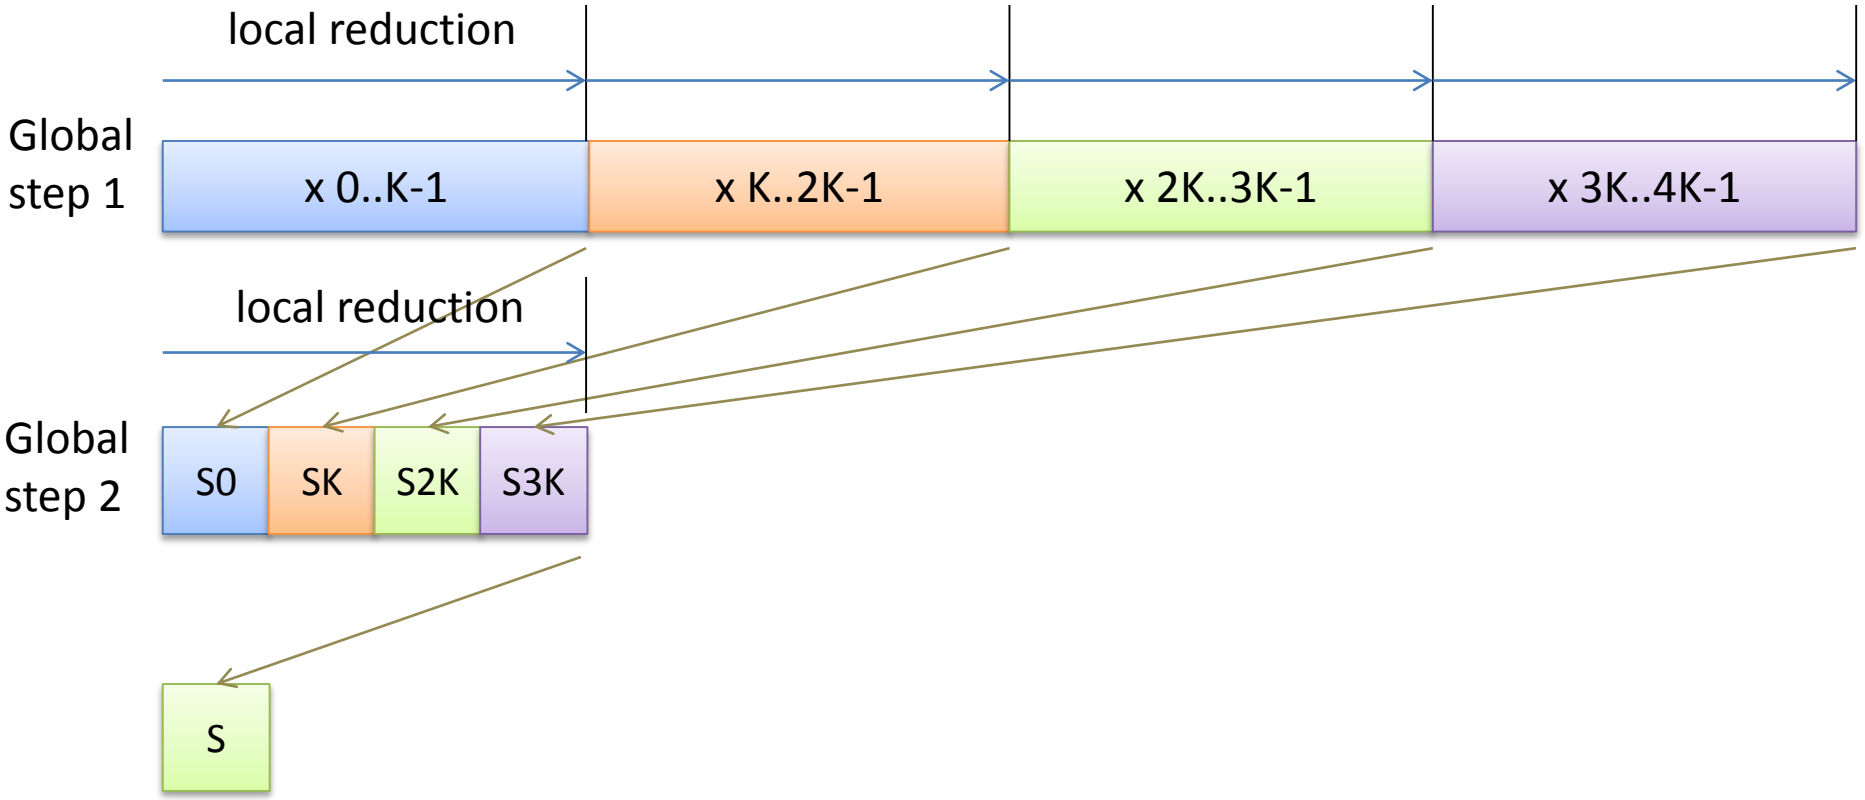

Increase per-thread work

- Small sequential sum in each thread
  - Further reduces global calls
  - More opportunities for latency hiding
- Start with a small loop
  - Careful with coalescence

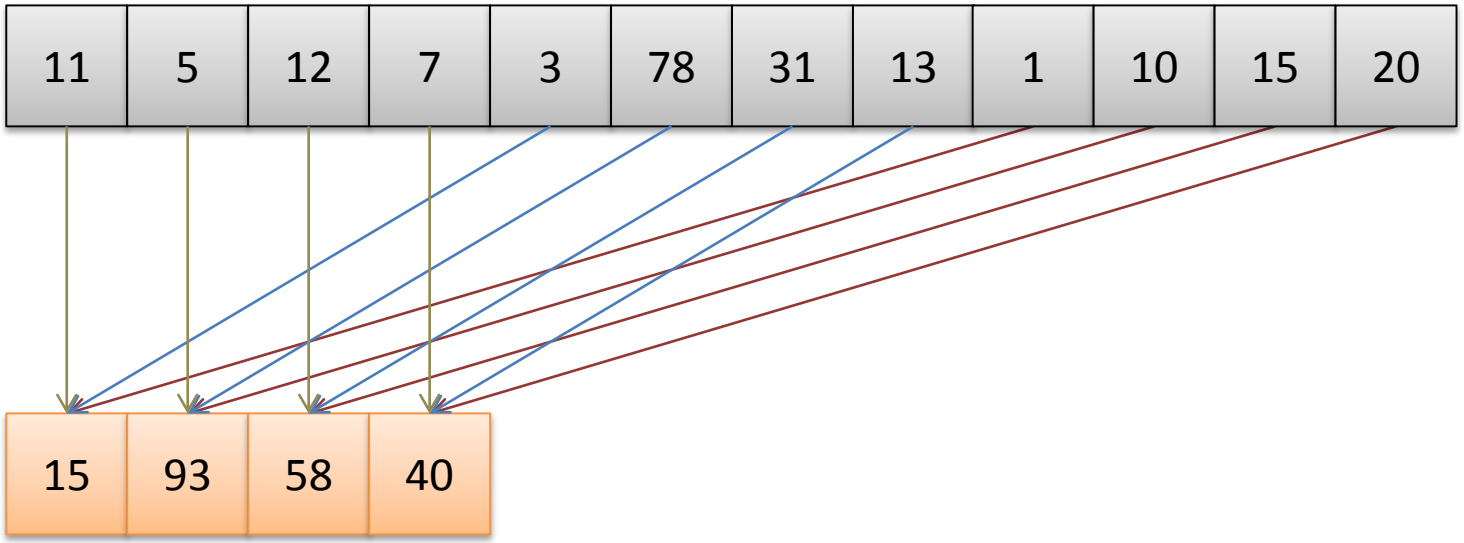

Parallel scan

Parallel scan

- *Scan* is also a very common operation
- Also known as prefix sum:

$$S[i] = \sum_{j<=i} T[j]$$

- Usually performed in-place

Sequential scan

- Very simple and efficient on a CPU:

```
for i : 1 .. N
        T[i] = T[i-1] + T[i]
```

- That's all!
- O(N) additions

Parallel scan

- Need to remove sequential dependency…

Parallel scan

- A first approach

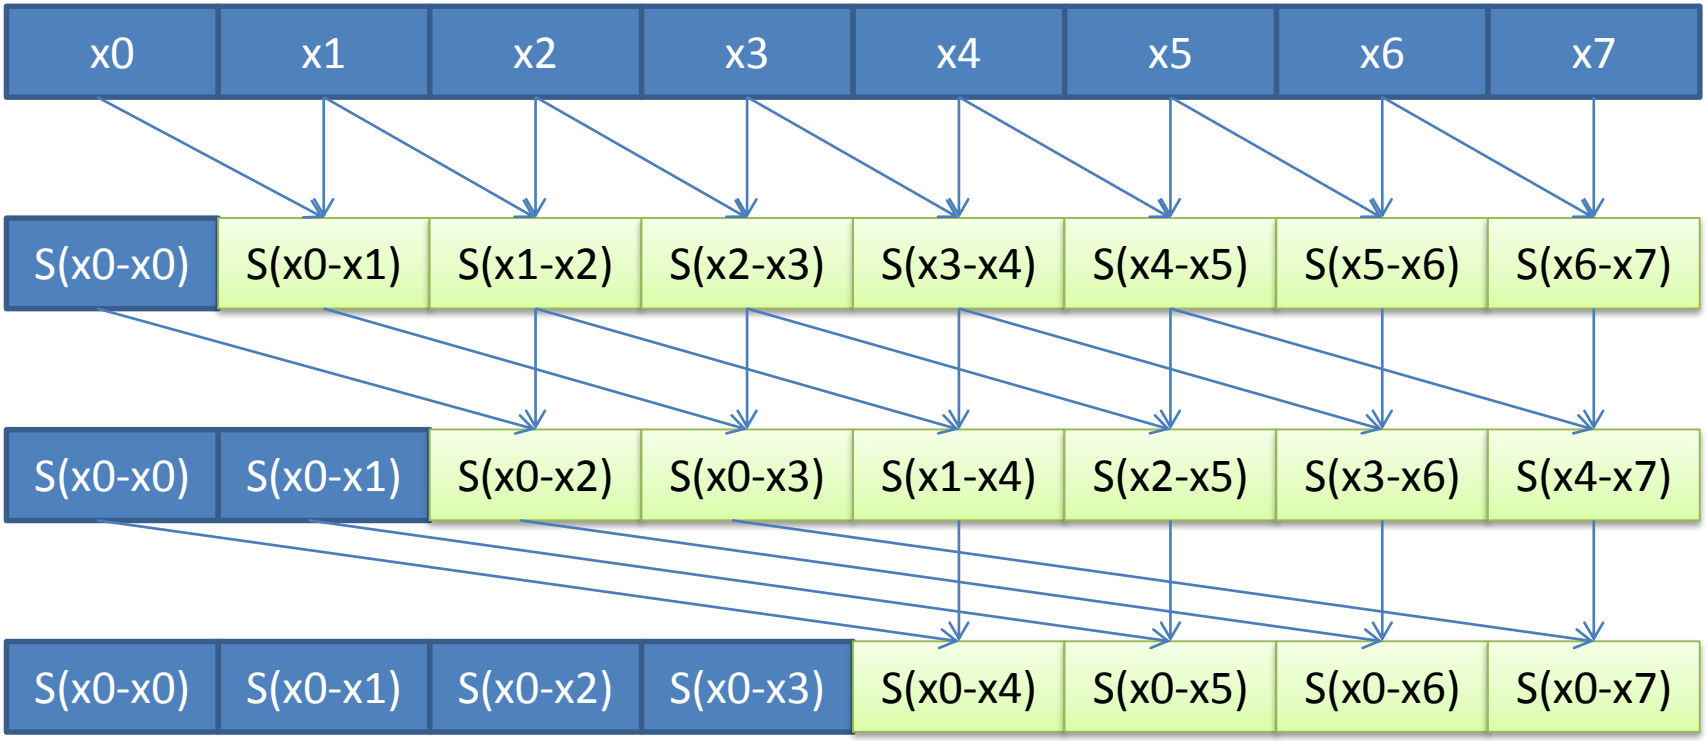

- Number of additions?

O(N log N)

Parallel scan

- First version:

  O(N log N) vs. O(N) in sequential

- Not *work-efficient*

Parallel scan (2)

- Main idea:

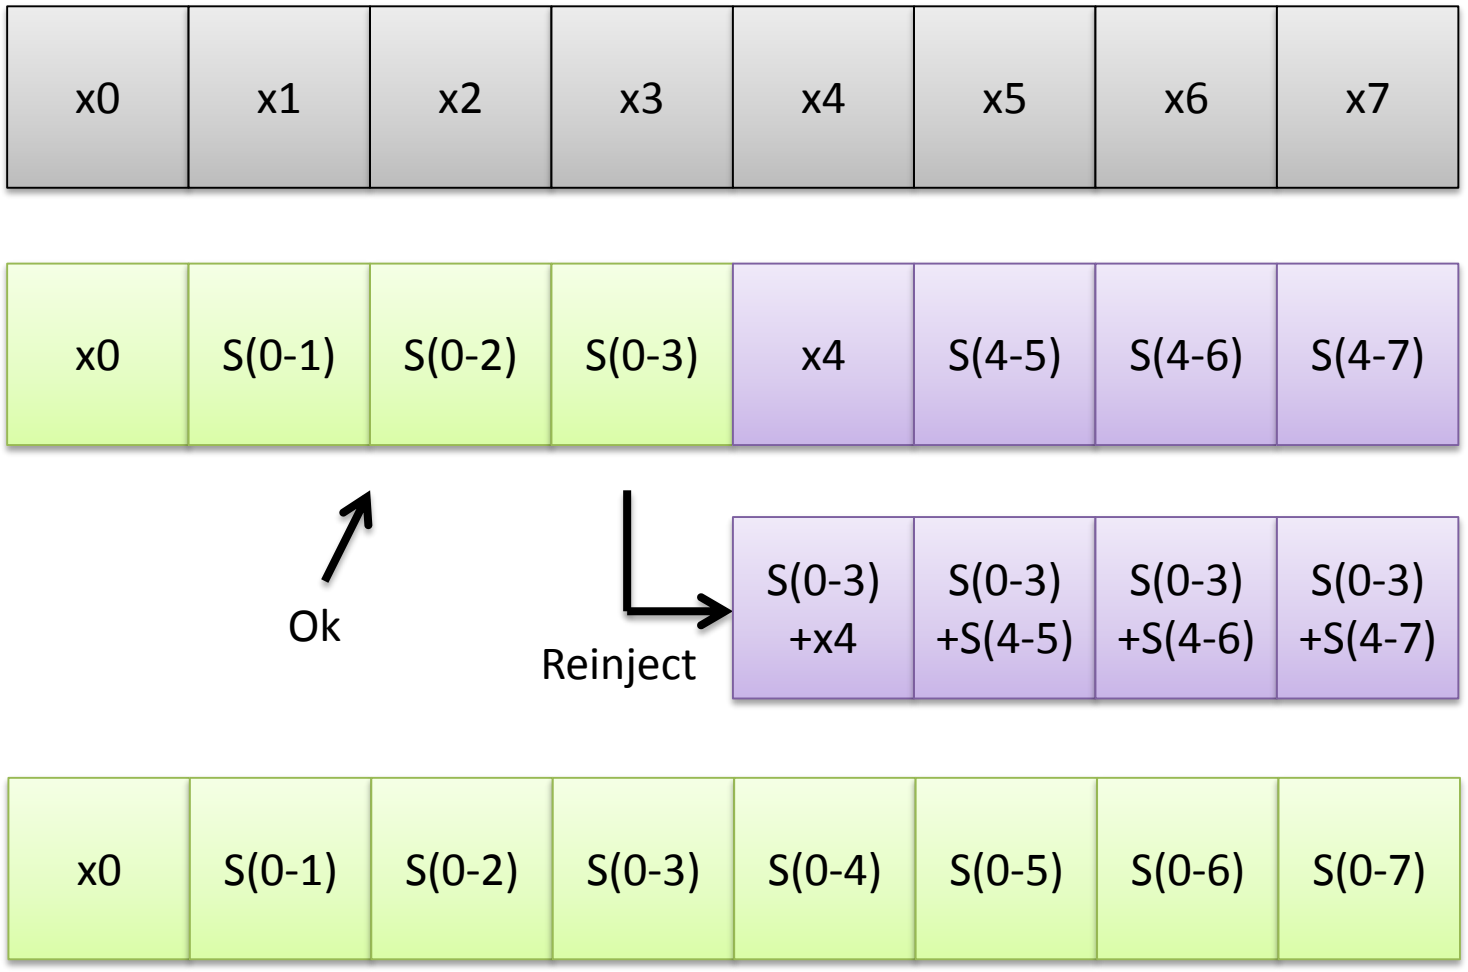

Parallel scan (2)

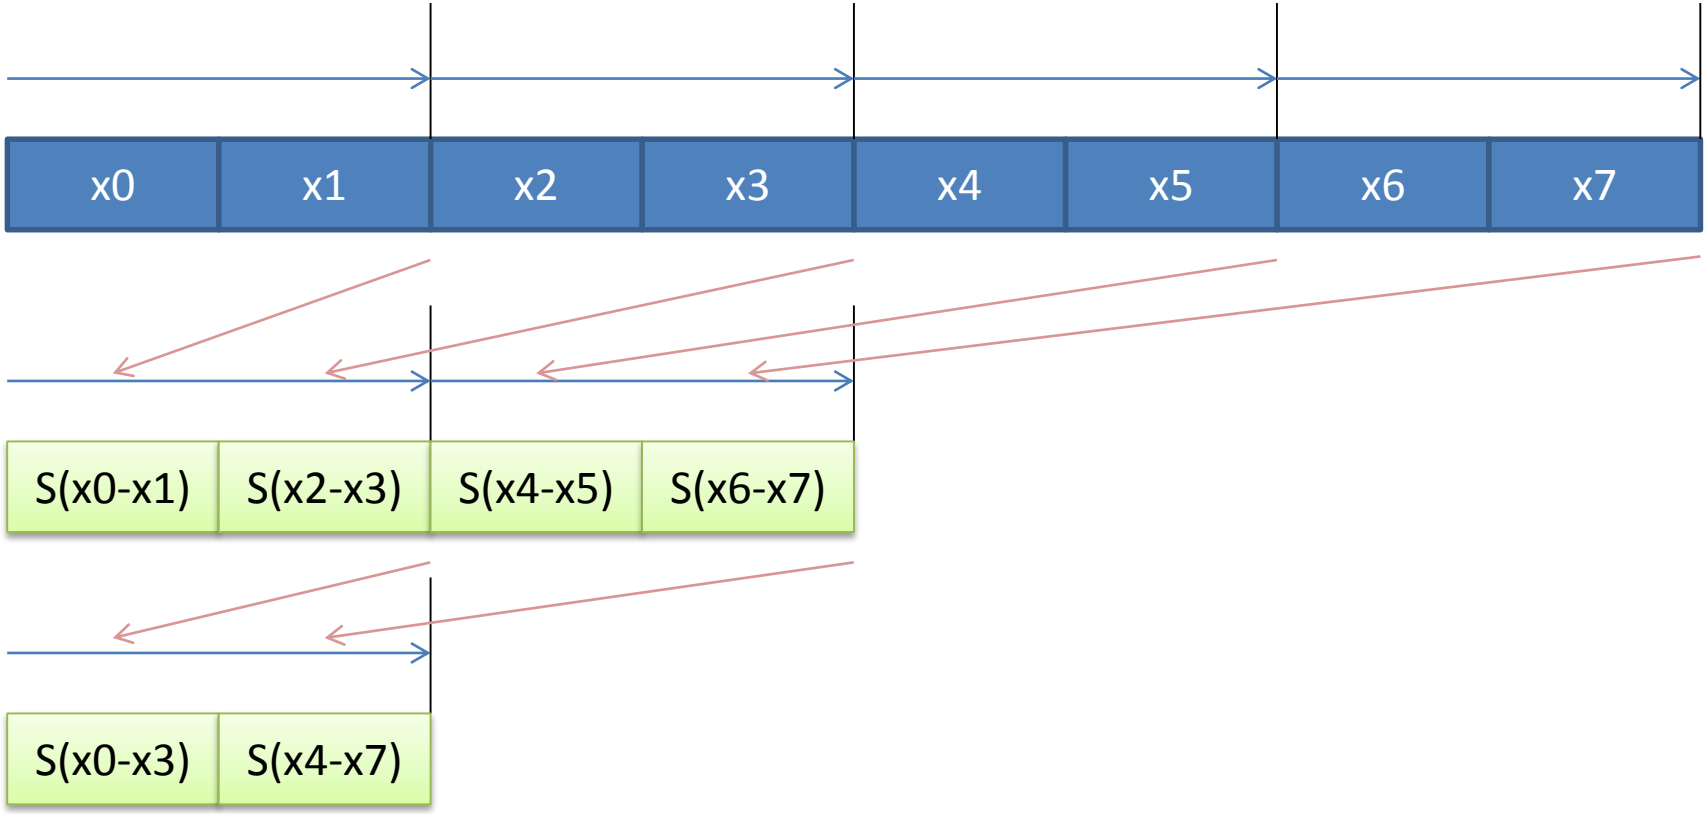

Parallel scan (2)

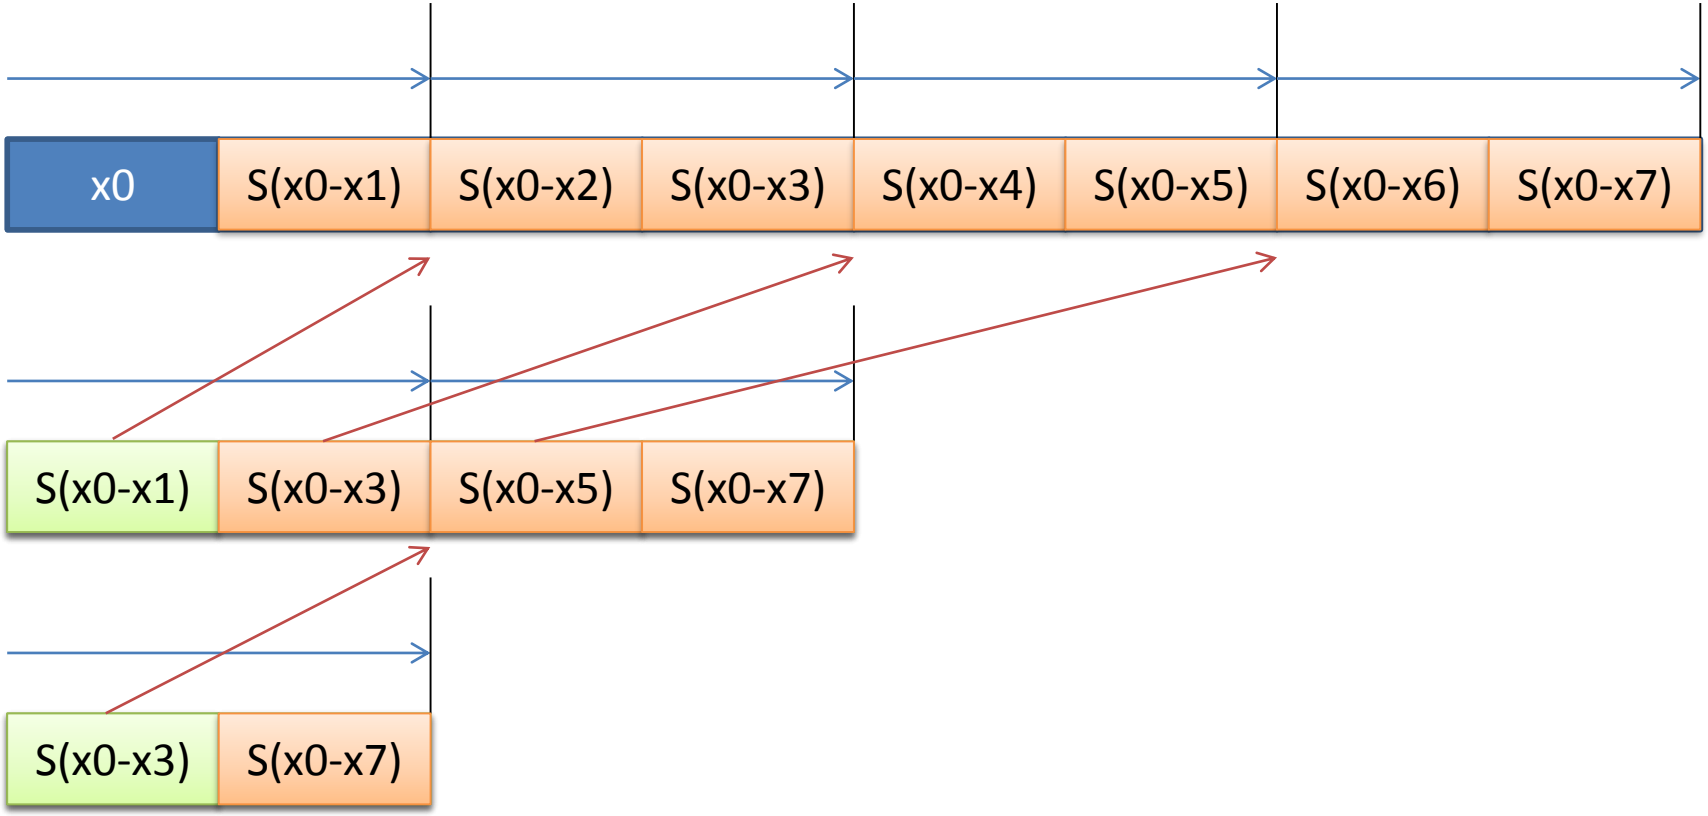

Parallel scan (2)

- Two phases:
  1. Perform a reduction (again!)

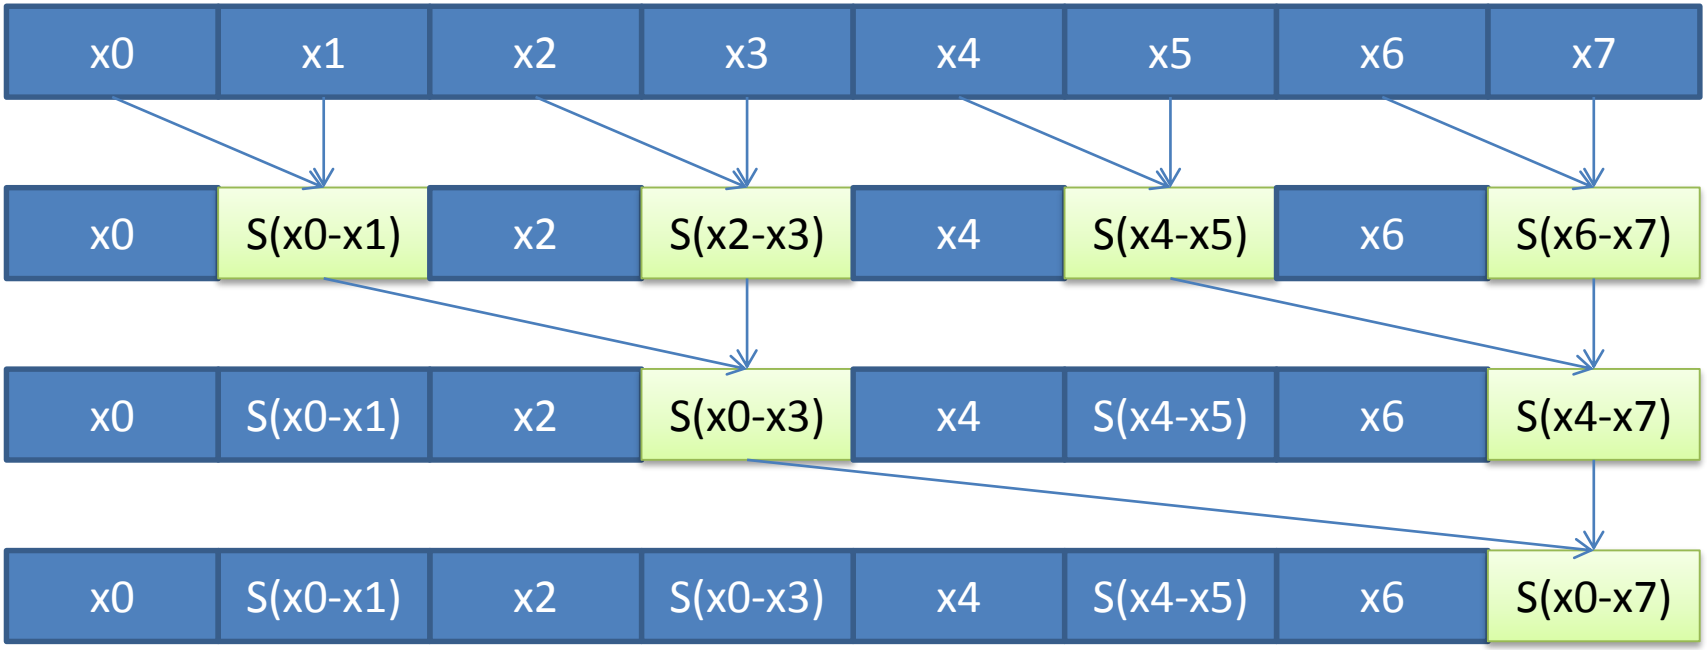

Parallel scan (2)

- Two phases:
  1. Perform a reduction (again!)
  2. Down-sweep

(slightly different result…)

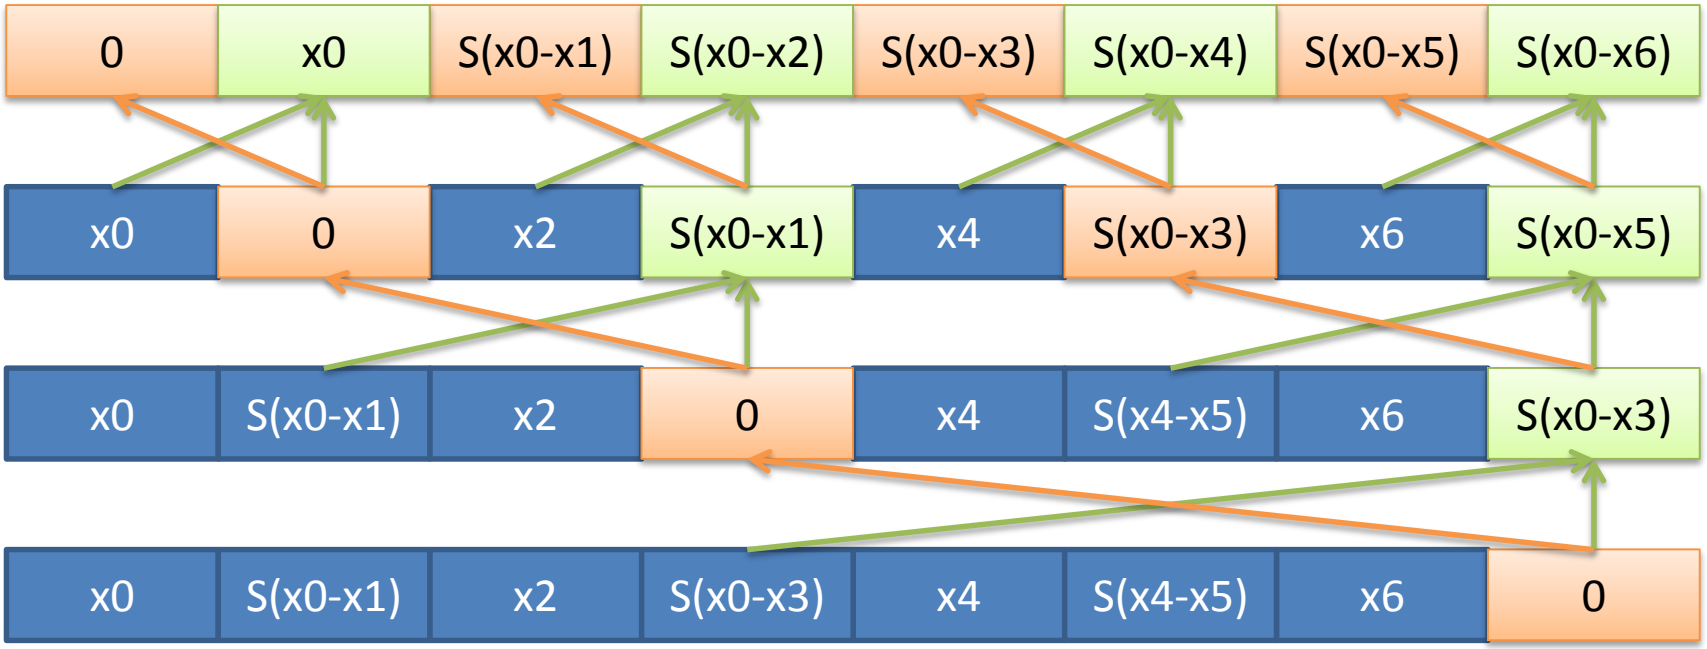

Parallel scan (2)

- Number of additions?

O(N)

- Memory access pattern?

Many conflicts!!

- Efficient implementation should consider this.

Home Assignment

- Parallel scan
  - Global memory, all array sizes
  - First version (slide 20)
  - Second version (slide 26)
  - Compare performance
  - Implement application 1 + versus atomicInc (bonus)
- Due date: 2013-04-19
- Return to: sylvain.lefebvre@inria.fr
  Subject: [OpenCL DM]

- Warning: This is a test

Application #1

- Sparse array compaction

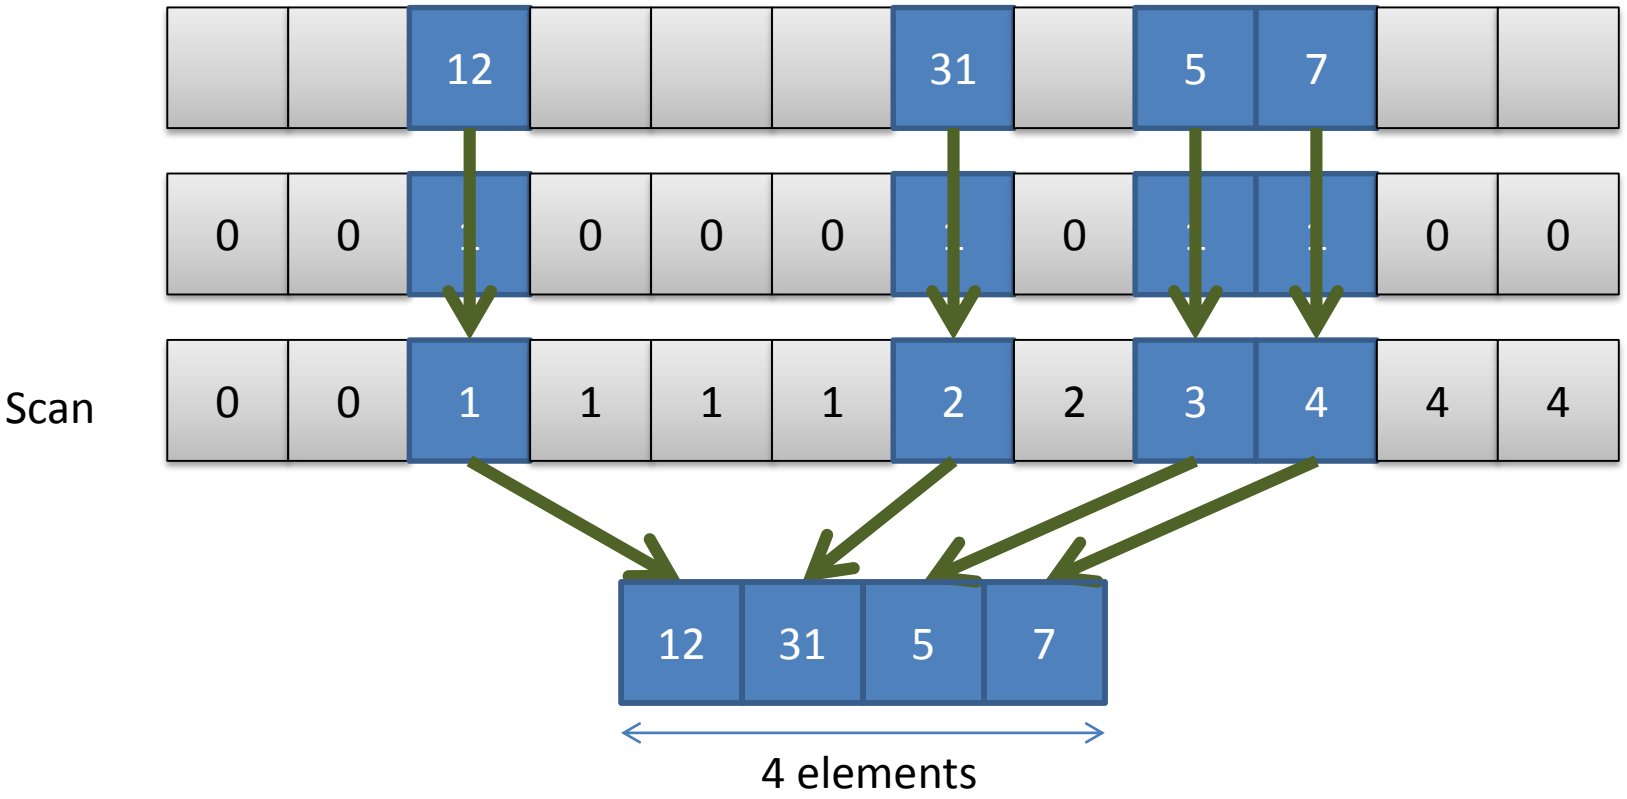

Application #2

- Special type of binary tree
  - Leaves at same level

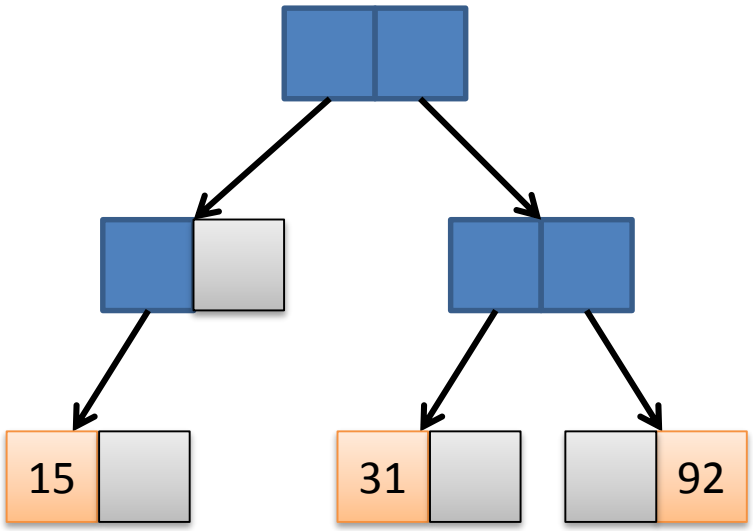

Binary-tree

- Captures a sparse array
  - Very common: Sparse linear algebra, integer sets, etc.
  - Fast retrieval O(log N)

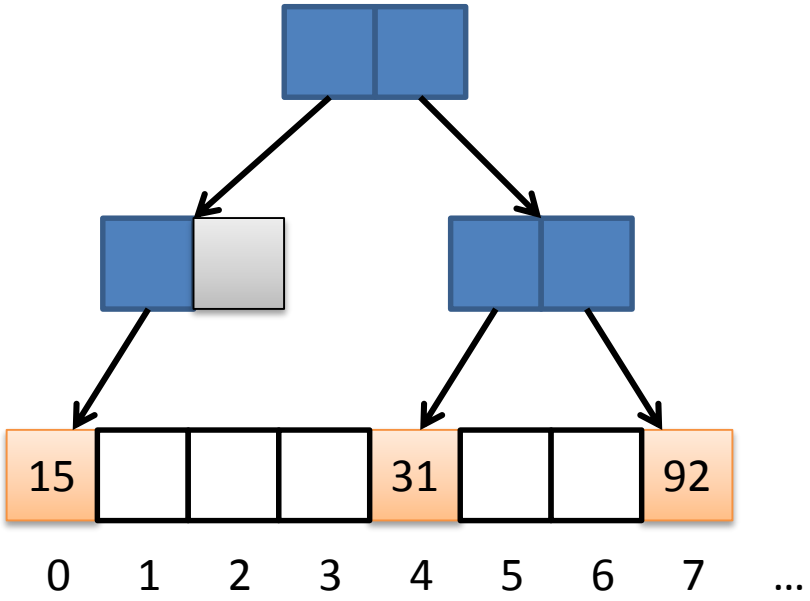

Encoding in memory

- All node records in a table
  - Internal nodes: Two integer indices (left / right child)
  - Leaves: Two integer data
- Called ‘autumnal tree’

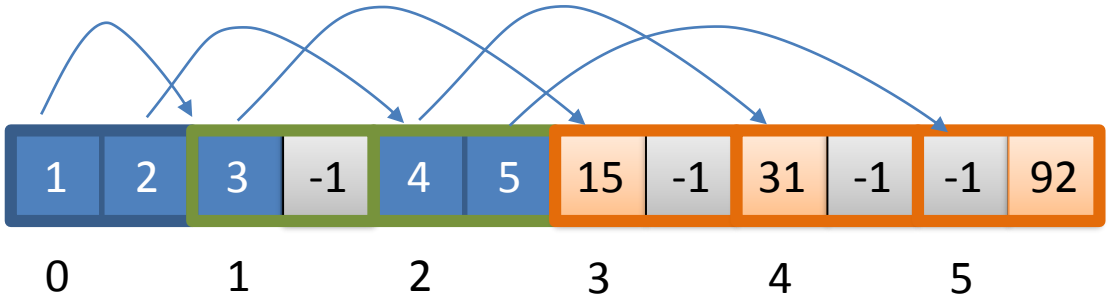

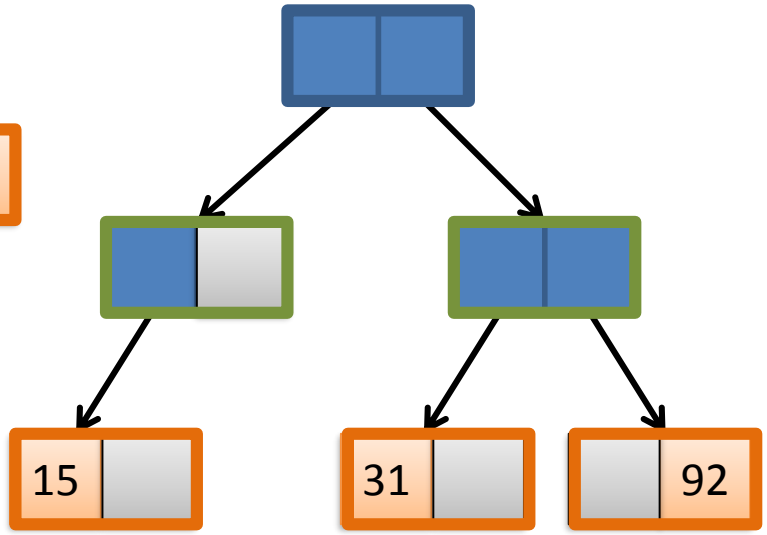

Building a binary-tree

- Input:

  Sparse table of integers

- Output:

  Node table

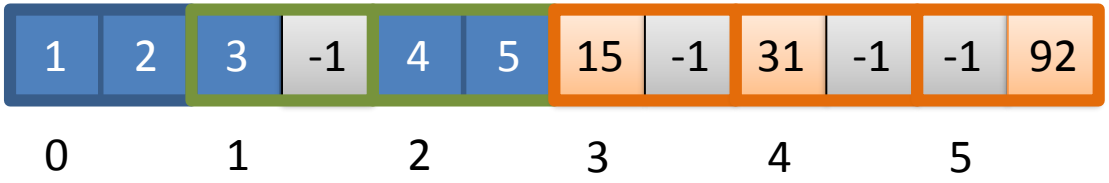

Principle

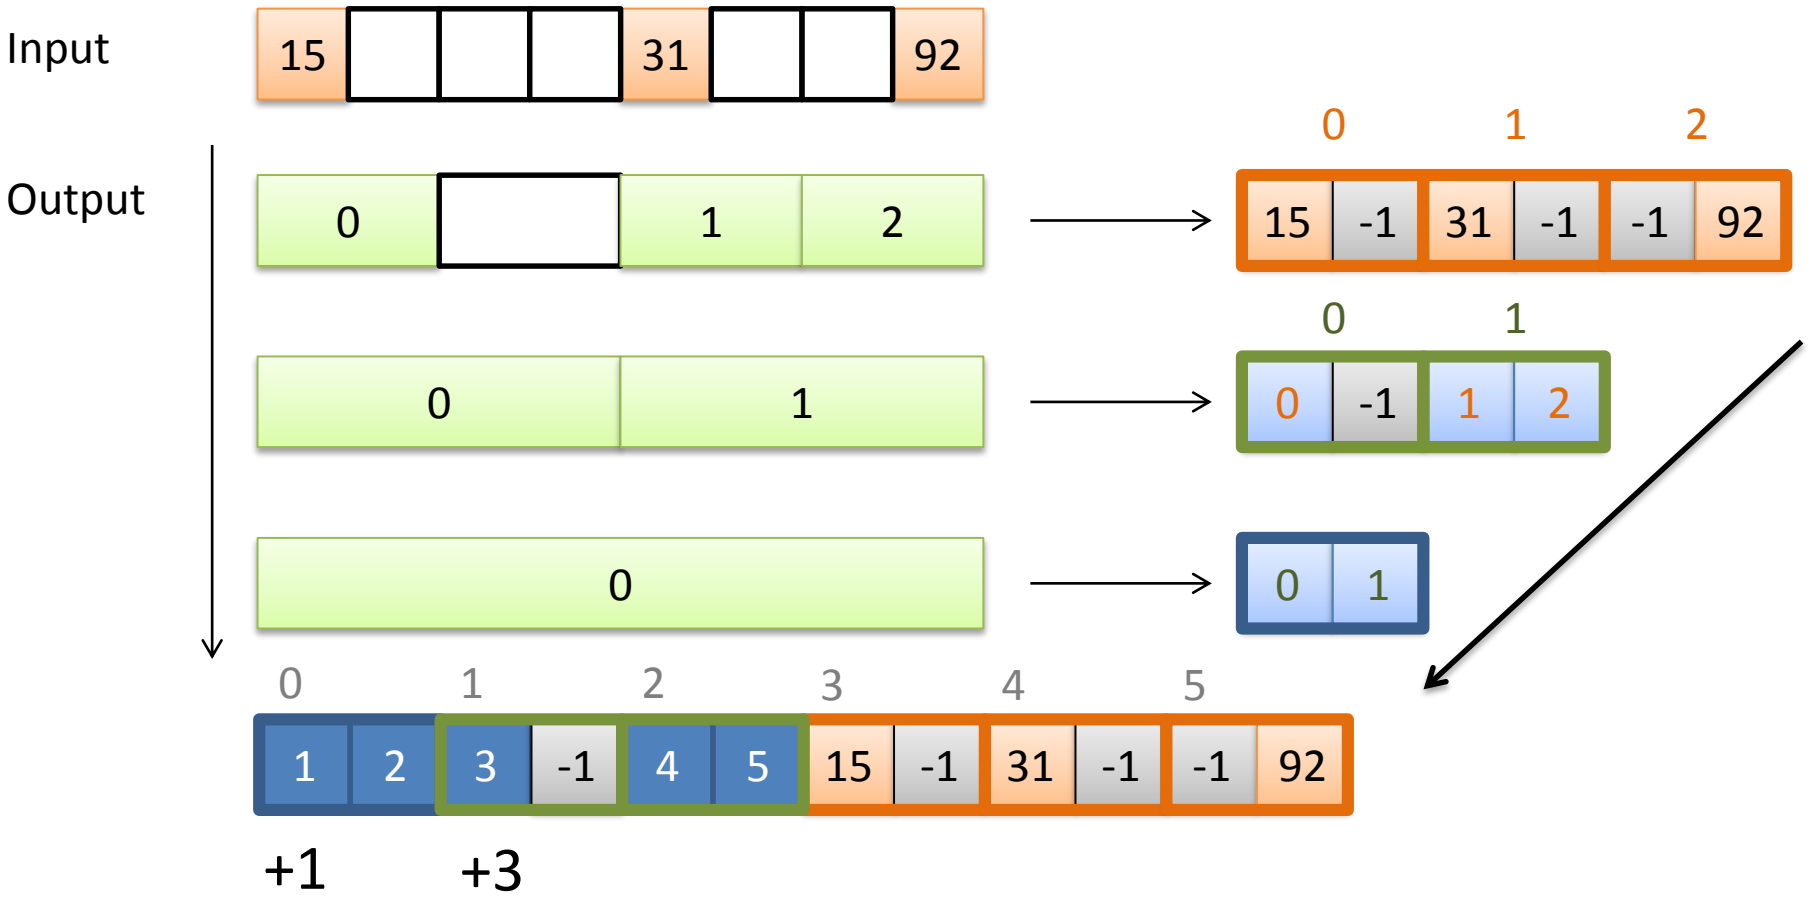
Input
15
31
92
Output
0
1
2
0
1
2
15
-1
31
-1
-1
92
0
1
0
-1
1
2
0
0
1
0
1
2
3
4
5
1
2
3
-1
4
5
15
-1
31
-1
-1
92
+1
+3

Let’s practice!

- Reduction
  - atomicAdd (TD)
  - Hierarchy, global
  - Hierarchy, global, better coalescence
  - Local – global approach

- Further reading

  M. Harris
  - Talk on Reduction
  - GPU Gems 3 *Parallel Prefix Sum (Scan) with CUDA*