

Introduction to data-parallelism and OpenCL

Sylvain Lefebvre - INRIA

Sorting

- Sort is everywhere
- Sequential sorts:
 - Selection sort $O(N^2)$
 - Heapsort $O(N \log N)$
 - Quicksort $O(N \log N)$
- Known best complexity: $O(N \log N)$

Parallel sort


- How can we sort in parallel?
 - For now, not concerned about performance
- Most sorting algorithms:
 - Number of comparisons depends on data
 - Example: Quicksort
 - Average $O(N \log N)$
 - Worst case $O(N^2)$
- Data dependency makes parallelism complex

Sorting networks

- **Data-independent** sorting algorithms
 - Whatever the data, the same operations occur
- There are several such algorithms:
 - Odd-even sort
 - Bitonic merge sort
 - ...

Odd-even sort

- Simple, first approach:



$O(N^2)$

Bitonic merge sort

- Better complexity: $O(N \log^2 N)$
- Based on ***bitonic*** sequences
 - Let's consider we only have values in $\{0,1\}$



bitonic



not bitonic



bitonic



not bitonic



bitonic






bitonic

Bitonic sequences have
at most 2 changes

Bitonic merge sort

- Bitonic sequences of numbers:



By rotation, two monotonic sub-sequences

0-1 principle


Theorem:

If a sorting network sorts every sequence of 0's and 1's, then it sorts every arbitrary sequence of values.

[Knuth 73]


Comparator

- Compares two halves of a bitonic 0-1 sequence




Comparator

- The two sub-sequences are bitonic
- The first is less-or-equal to the second




Bitonic merge



Bitonic sort

- Divide and conquer approach (recursive)
 - Two operations: **Sort** and **Merge**



Implementation

- Trivial cases:
 - Sort and Merge on $N = 1$
- Recursion (CPU):

```
Sort( T, L,R, d) {  
    if ( R-L+1 > 1 ) {  
        Sort( T, L, (R+L)/2,    ASC  )  
        Sort( T, (R+L)/2+1,R,  DESC )  
        Merge( T , L,R, d)  
    }  
}
```

```
Merge( T, L,R, d) {  
    if ( R-L+1 > 1 ) {  
        Compare( T, L,R ,d )  
        Merge( T , L, (L+R)/2    ,d )  
        Merge( T , (L+R)/2+1,R ,d )  
    }  
}
```

Implementation


```
Compare( T, R,L, d )
    k = (R-L+1)/2
    for (i = 0 ; i < k ; i++) {
        if (d == ASC) {
            if ( T[L+i] < T[L+i+k] )
                swap(T,L+i,L+i+k)
        } else {
            if ( T[L+i] > T[L+i+k] )
                swap(T,L+i,L+i+k)
        }
    }
}
```

```
Sort( T, L,R, d) {
    if ( R-L+1 > 1 ) {
        Sort( T, L, (R+L)/2, ASC )
        Sort( T, (R+L)/2+1,R, DESC )
        Merge( T , L,R, d)
    }
}
```

```
Merge( T, L,R, d) {
    if ( R-L+1 > 1 ) {
        Compare( T, L,R ,d )
        Merge( T , L, (L+R)/2 ,d )
        Merge( T , (L+R)/2+1,R ,d )
    }
}
```


Implementation in OpenCL

- Consider global comparison pattern



Implementation in OpenCL

- $\log(N)$ stages



Number of comparisons

- $\log(N)$ stages
- Stage i has i columns
- Number of columns:

$$\text{Sum}(i=0 \dots \log N) (i) = \log N (\log N + 1) / 2$$

- Number of comparisons:
 $N/2 * (\log N (\log N + 1) / 2) \rightarrow O(N \log^2 N)$

Let's practice

- Odd-even sort
 - Global memory
- Bitonic sort
 - First, recursive on CPU
 - Second, on the GPU (global memory)

Credits

- http://http.developer.nvidia.com/GPUGems2/gpugems2_chapter46.html
- <http://www.inf.fh-flensburg.de/lang/algorithmen/sortieren/bitonic/bitonicen.htm>